

Adwait Jog

111N IST Building
Penn State University
University Park, PA, 16802
Homepage: <http://adwaitjog.github.io>

Email: adwait@cse.psu.edu
Alt. Email: adwaitjog@gmail.com
Skype ID: adwaitjog

RESEARCH INTERESTS

My research interests lie broadly in the area of computer architecture and systems.

PROFESSIONAL EXPERIENCE

- **The Pennsylvania State University, Ph.D. Candidate** **Fall 2009 - Summer 2015**
Advisor: Prof. Chita Das, High Performance Computing Lab (HPCL) **University Park, PA**
 - Proposed techniques for efficient execution of multiple applications on next generation GPUs.
 - Proposed criticality-aware memory system for GPUs.
 - Proposed a coordinated scheduling and prefetching mechanism to improve GPU performance.
 - Proposed warp scheduling policies to mitigate contention in GPU memory system.
 - Traded-off non-volatility of STT-RAM for lower write latency and energy.
- **NVIDIA Research, Graduate Research Intern** **Summer 2013**
Manager: Steve Keckler **Santa Clara, CA**
Mentors: Evgeny Bolotin, Zvika Guz, Mike Parker
Researched on efficient execution of multiple contexts/applications on next generation GPUs. The results of this work are published in GPGPU 2014 (co-located with ASPLOS 2014).
- **Intel Labs, Graduate Research Intern** **Summer 2012**
Managers: Srihari Makineni, Ravi Iyer **Hillsboro, OR**
Mentors: Xiaowei Jiang, Li Zhao
Implemented and evaluated micro-architecture techniques for Intel's ultra-low power core (*Siskiyou*). This infrastructure is released to universities to perform research on energy-efficient architectures.
- **Intel Corp., Graduate Technical Intern** **Summer 2011**
Manager: Sridhar Lakshmanmoorthy **Hillsboro, OR**
Mentor: Ramadass Nagarajan
Performed detailed studies and proposed techniques for designing a QoS aware interconnect fabric for the Intel's next generation heterogeneous SoCs.

EDUCATION

The Pennsylvania State University, University Park
Ph.D. Candidate in Computer Science and Engineering
Advisors: Chita Das and Mahmut Kandemir

Fall 2009 - Summer 2015

National Institute of Technology, Rourkela, India
Bachelor of Technology (*Hons.*) in Electronics and Instrumentation Engineering

Fall 2005 - Spring 2009

AWARDS and HONORS

- Outstanding Graduate Research Assistant Award, CSE, Penn State, 2014
- NVIDIA Graduate Fellowship 2013, Finalist

- Best Paper Nomination, PACT 2013
- Student Travel Grants: MICRO 2014, ASPLOS 2014, ISCA 2013, ASPLOS 2013
- College of Engineering Fellowship, Penn State University, 2009
- Summer Research Fellowship, School of Computing, National University of Singapore, 2008
- Summer Research Fellowship, Indian Academy of Sciences, 2007
- Undergraduate Scholarship for being in Top 1% in All India Engineering Entrance Exam, 2005

PUBLICATIONS

- [**ISCA 2015**] Nandita Vijaykumar, Gennady Pekhimenko, **Adwait Jog**, Abhishek Bhowmick, Rachata Ausavarungnirun, Onur Mutlu, Chita Das, Mahmut Kandemir, Todd Mowry, *A Case for Core-Assisted Bottleneck Acceleration in GPUs: Enabling Efficient Data Compression*, In 42nd International Symposium on Computer Architecture (ISCA), Portland, OR, June, 2015
- [**MICRO 2014**] Onur Kayiran, Nachiappan CN, **Adwait Jog**, Rachata Ausavarungnirun, Mahmut Kandemir, Gabriel Loh, Onur Mutlu, Chita Das, *Managing GPU Concurrency in Heterogeneous CPU-GPU Architectures*, In 47th International Symposium on Micro Architecture (MICRO), Cambridge, UK, Dec 2014
- [**GPGPU 2014**] **Adwait Jog**, Evgeny Bolotin, Zvika Guz, Mike Parker, Steve Keckler, Mahmut Kandemir, Chita Das, *Application-aware Memory System for Fair and Efficient Execution of Concurrent GPGPU Applications*, In General-Purpose Computation on Graphics Processing Unit (GPGPU-7), co-located with ASPLOS, Salt Lake City, UT, USA, March, 2014
- [**PACT 2014**] Wei Ding, Mahmut Kandemir, Diana Guttman, **Adwait Jog**, Chita Das, Praveen Yedlapalli, *Trading Cache Hit Rate for Memory Performance*, In 23rd International Conference on Parallel Architectures and Compilation Techniques (PACT), Edmonton, Canada, August, 2014
- [**ASPLOS 2013**] **Adwait Jog**, Onur Kayiran, Nachiappan CN, Asit Mishra, Mahmut Kandemir, Onur Mutlu, Ravi Iyer, Chita Das, *OWL: Cooperative Thread Array Aware Scheduling Techniques for Improving GPGPU performance*, In 18th International Conference on Architectural Support for Programming Languages and Operating Systems (ASPLOS), Houston, TX, USA, March, 2013
- [**ISCA 2013**] **Adwait Jog**, Onur Kayiran, Asit Mishra, Mahmut Kandemir, Onur Mutlu, Ravi Iyer, Chita Das, *Orchestrated Scheduling and Prefetching for GPGPUs*, In 40th International Symposium on Computer Architecture (ISCA), Tel Aviv, Israel, June, 2013
- [**PACT 2013, Best Paper Nominee**] Onur Kayiran, **Adwait Jog**, Mahmut Kandemir, Chita Das, *Neither More Nor Less: Optimizing Thread-Level Parallelism for GPGPUs*, In 22nd International Conference on Parallel Architectures and Compilation Techniques (PACT), Edinburgh, Scotland, September, 2013
- [**DAC 2012**] **Adwait Jog**, Asit Mishra, Cong Xu, Yuan Xie, Vijaykrishnan Narayanan, Ravi Iyer, Chita Das, *Cache Revive: Architecting Volatile STT-RAM Caches for Enhanced Performance in CMPs*, In 49th Design Automation Conference (DAC), San Francisco, USA, June 2012

TEACHING EXPERIENCE

- **Co-Instructor**, CMPEN 331, Computer Organization and Design **Fall 2014**
- **Co-Instructor**, CMPEN 331, Computer Organization and Design **Spring 2014**
- **Teaching Assistant**, CMPEN 431, Introduction to Computer Architecture **Spring 2010**
- **Teaching Assistant**, CMPEN 471, Logic Design of Digital Systems **Fall 2009**

TALKS AND POSTER SESSIONS

- The Future of Parallel Computing with GPUs
 - The College of William and Mary, Feb 2015
 - University of Utah, Mar 2015
 - Temple University, Mar 2015
 - AMD Research, Mar 2015
 - UC Riverside, Apr 2015
 - Intel Labs, Apr 2015
- Application-aware Memory System for Fair and Efficient Execution of Concurrent GPU Applications,
 - GPGPU-7 Workshop (co-located with ASPLOS 2014), Salt Lake City, UT, March 2014
 - Intern Talk, NVIDIA Research, Santa Clara, CA, Sept 2013
- Mitigating and Masking the Limitations of GPU Memory Systems,
 - Intern Talk, NVIDIA Research, Santa Clara, CA, June 2013
- Orchestrated Scheduling and Prefetching for GPGPUs,
 - ISCA 2013, Tel Aviv, Israel, June 2013
- OWL: Cooperative Thread Array Aware Scheduling Techniques for Improving GPGPU performance,
 - ASPLOS 2013, Houston, TX, March 2013
- Cache Revive: Architecting Volatile STT-RAM Caches for Enhanced Performance in CMPs,
 - DAC 2012, San Francisco, CA, June 2012
 - Poster presentation at IUCRC NEXYS Workshop, Pittsburgh, PA

PROFESSIONAL SERVICE AND MEMBERSHIPS

- Proceedings and Submission Chair, ANCS 2015
- Reviewer (Conferences): I was on External Review Committee for DAC 2013, HPCA 2013, MICRO 2012, and ICCD (2014, 2013)
- On-Behalf Reviewer (Conferences): ASPLOS, ISCA, HPCA, MICRO, DATE, DAC, PACT, IPDPS
- Reviewer (Journals): IEEE Transactions on Computers (TC), ACM Transactions on Embedded Computing (TECS), ACM Transactions on Design Automation of Electronic Systems (TODAES), IEEE Journal on Computer Architecture Letters (CAL)
- Student Member of ACM, IEEE, ACM SIGARCH